



(a) Training phase - shared learner

Each bus contributes a local transition and per-agent reward.



(b) Execution phase - decentralized

Frozen shared weights; each bus acts on its own local observation.



No centralized real-time action selector
and no weight updates at deployment